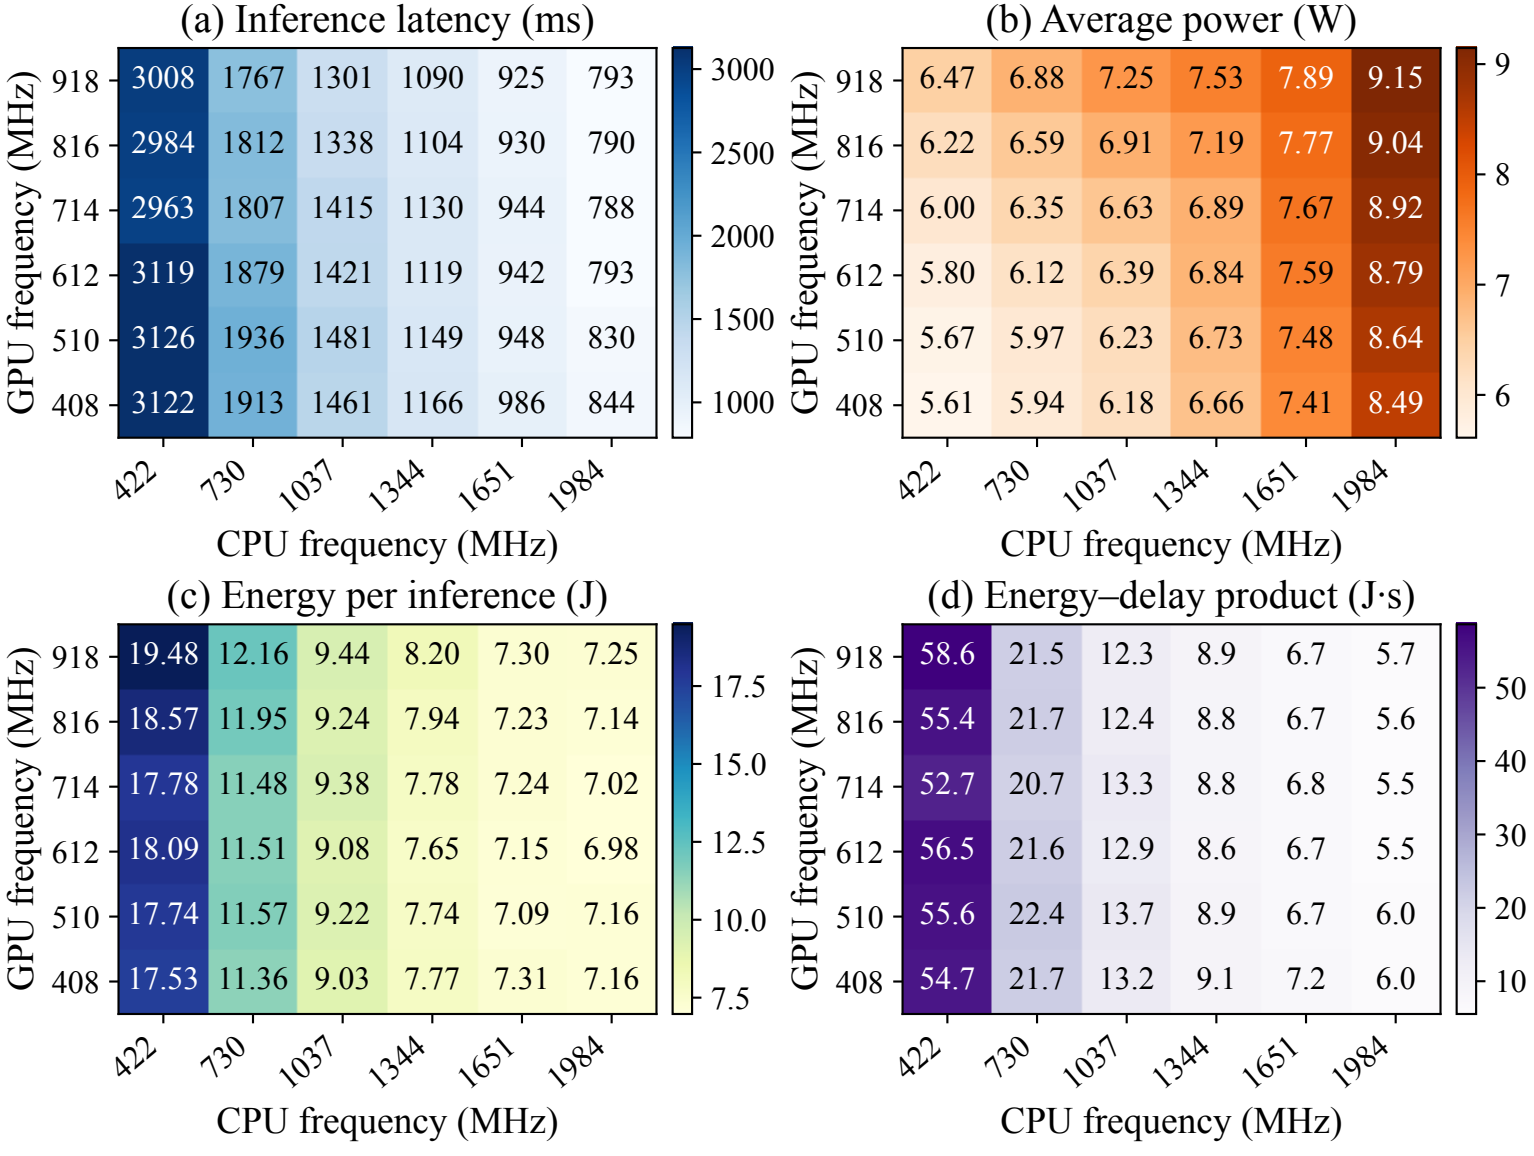
(a) Inference latency (ms)
GPU frequency (MHz)
CPU frequency (MHz)
918 816 714 612 510 408
422 730 1037 1344 1651 1984
3008 1767 1301 1090 925 793
2984 1812 1338 1104 930 790
2963 1807 1415 1130 944 788
3119 1879 1421 1119 942 793
3126 1936 1481 1149 948 830
3122 1913 1461 1166 986 844
3000 2500 2000 1500 1000
(b) Average power (W)
GPU frequency (MHz)
CPU frequency (MHz)
6.47 6.88 7.25 7.53 7.89 9.15
6.22 6.59 6.91 7.19 7.77 9.04
6.00 6.35 6.63 6.89 7.67 8.92
5.80 6.12 6.39 6.84 7.59 8.79
5.67 5.97 6.23 6.73 7.48 8.64
5.61 5.94 6.18 6.66 7.41 8.49
9 8 7 6
(c) Energy per inference (J)
GPU frequency (MHz)
CPU frequency (MHz)
19.48 12.16 9.44 8.20 7.30 7.25
18.57 11.95 9.24 7.94 7.23 7.14
17.78 11.48 9.38 7.78 7.24 7.02
18.09 11.51 9.08 7.65 7.15 6.98
17.74 11.57 9.22 7.74 7.09 7.16
17.53 11.36 9.03 7.77 7.31 7.16
17.5 15.0 12.5 10.0 7.5
(d) Energy–delay product (J·s)
GPU frequency (MHz)
CPU frequency (MHz)
58.6 21.5 12.3 8.9 6.7 5.7
55.4 21.7 12.4 8.8 6.7 5.6
52.7 20.7 13.3 8.8 6.8 5.5
56.5 21.6 12.9 8.6 6.7 5.5
55.6 22.4 13.7 8.9 6.7 6.0
54.7 21.7 13.2 9.1 7.2 6.0
50 40 30 20 10